

Started on	Saturday, 27 July 2024, 10:25 PM
State	Finished
Completed on	Saturday, 3 August 2024, 2:30 PM
Time taken	6 days 16 hours
Marks	43/43
Grade	10 out of 10 (100%)

Question 1

Correct

Mark 1 out of 1

Decision trees are machine learnable models that learn a mapping from an input space to an output space based on a finite amount of training samples exemplifying this mapping. Concretely, a decision tree model is

- ☐ a. a special kind of deep neural network
- ☒ b. defining an axis-aligned decision boundary
- ☐ c. restricted to categorical output features
- ☐ d. restricted to categorical inputs
- ☒ e. a directed tree of conditions as nodes and discrete outputs as leaves
- ☒ f. a partition of the input space

Your answer is correct.

Correct

Marks for this submission: 1/1.

Question 2

Correct

Mark 5 out of 5

A decision tree consists of

- connection of a superset to its subsets which are obtained via splitting according to a splitting condition
- subset of the input space together with a hierarchy of conditions to further break it down
- subsets of the input space that are not further split; these can be associated with an output class label
- subsets of the input space together with a splitting condition that uses one feature
- the first splitting on the complete input space, called

- edges
- subtree
- leaf nodes
- internal nodes
- root node

Your answer is correct.

Correct

Marks for this submission: 5/5.

Question 3

Correct

Mark 1 out of 1

A decision tree is built by iterative splitting of nodes. What are typical stop criteria for quitting the splitting at a node?

- ☒ a. low number of training samples in the node
- ☐ b. further splitting would exceed a maximum width
- ☐ c. low variance reduction for any possible split
- ☒ d. further splitting would exceed a maximum depth
- ☐ e. low information gain for any possible split
- ☐ f. low Gini impurity of the node

Your answer is correct.

Correct

Marks for this submission: 1/1.

Question 4

Correct

Mark 4 out of 4

To decide whether a proposed split is better than another, decision tree learning algorithms utilize one of several methods to quantify "purity" of a subset with respect to the output classes/values.

Let's first **recap the formulas**:

For output variable Y , feature X , and set to split S , what is the score $C(X; Y)$ for improvement of purity when splitting at X ?

- : $C(X = x) = 1 - \sum_y P(y|x)^2$; $C(X; Y) = C(X = \text{any}) - \sum_x P(X = x)C(X = x) = -\sum_y P(y)^2 + \sum_{y,x} P(x)P(y|x)^2$
- : $C(X; Y) = H(Y) - H(Y|X) = -\sum_y P(y) \log(P(y)) + \sum_{y,x} P(x)P(y|x) \log(P(y|x))$
- : $C(X; Y) = \frac{1}{|S|^2} \left(\sum_{s,s' \in S} k(y_s, y_{s'}) - \sum_x \sum_{s,s' \in S_{X_i=x}} k(y_s, y_{s'}) \right)$

Note: Notice the similarity of the first two! In essence, they both sum some strengthened form of the conditional probabilities; one $P(y|x)P(y|x)$, the other $P(y|x)\log(P(y|x))$. Thus, there is not much of a difference to be expected in the outcomes when using either of the two.

Now we can have a closer look at **their properties**.

Which of the following holds true about the purity measures that we have seen in the lecture?

- works both for regression and classification tasks; works for regression tasks, and the rest for classification.
- measures the probability of correctly guessing the class of a sample drawn from the split.
- measures how much the training sample output values within each split deviate from one another.
- measures how far the respective distributions are from being random.
- Entropy lies at the heart of : It measures the change in entropy when splitting.
- was the first used metric, while was first introduced in CART for classification trees.

Your answer is correct.

Correct

Marks for this submission: 4/4.

Question 5

Correct

Mark 4 out of 4

Let's recapitulate the decision tree learning algorithm step by step, and start with the overview here.

In the following pseudo-code, insert the available comments at the correct position which describe the rational of the respective steps.

```
Given: input space S, training samples t, validation samples v, features X, output feature Y
Goal: tree consisting of
* inner nodes Node(feature, input space region),
* leaves Leaf(input space region), and
* connecting edges represented by tests Test("point[{some feature}] </>= {some value of feature}"),
where each node defines a subtree.

# start at the complete input space
tree <- Leaf(S)
# 1. 
FOR leaf in tree.leaves:
    subtree <- split_leaf(leaf, t, X, Y)
    replace leaf in tree by subtree
    tree.leaves.append( leaves in splitting_feature.split)

# 2. 
FOR leaf in tree.leaves:
    # 
    leaf.output <- determine_output(leaf.samples)

# 3. 
all_subtrees <- sort(nodes in tree) # go through subtrees bottom-up or top-down
FOR subtree in subtrees:
    replace subtree in tree by prune_subtree(subtree, v)
```

Your answer is correct.

Correct

Marks for this submission: 4/4.

Question 6

Correct

Mark 4 out of 4

Let's take a closer look at the iterative splitting of the tree nodes. Fill in the comments at the right position.

```
FUNCTION split_Leaf(Leaf leaf, TrainingData t, InputFeatures X, OutputFeature Y, Bool is_part_of_random_forest = true):  
  # 1.   
  IF maximum depth reached or |leaf| < minimum samples:  
    return leaf # no splitting  
  
  X' <- X  
  # 2.   
  IF is_part_of_random_forest:  
    X' <- randomly sample from features X  
  
  # 3.   
  leaf.samples <- points in t that lie in leaf.subset # if not already defined  
  FOR feature F in X':  
    F.split <- determine_subtree_for_feature(leaf, F)  
    #   
    F.improvement <- calculate_purity_improvement_on_training_set_Y_values_if_applying_F_split  
  
  # 4.   
  splitting_feature <- argmax_F F.improvement  
  RETURN splitting_feature  
END FUNCTION
```

```
FUNCTION determine_subtree_for_feature(Leaf leaf, feature F):  
  # in the categorical case:   
  IF F categorical:  
    edges = [ Test("point[{F}] == {value}") for value in F.values ]  
    split <- Node(F, leaf.samples) connected via edge to Leaf({p in leaf | p passes edge.test}) for edge in edges  
  
  # in the numerical case:   
  ELSEIF F numerical:  
    # 1.   
    F.possible_splitting_points <- F-values halfway between points  
  
    # optional:   
    F.possible_splitting_points <- splitting_points_pruned_to_define_more_consistent_intervals  
  
    # 2.   
    FOR split_point in F.possible_splitting_points:  
      edges = [ Test("point[{F}] <= {split_point}"), Test("point[{F}] > {split_point}") ]  
      F.possible_splits.append( Node(F) connected via edge to Leaf({p in leaf | p passes edge.test}) for edge in edges )  
  
    # select best splitting point  
    split <- split_subtree_in_F_possible_splits_with_maximum_quality_gain  
  
  RETURN split  
END FUNCTION
```

```
FUNCTION determine_output(TrainingData training_data_subset):  
  IF Y categorical:  
    output <- most_common_Y_value_of_points_in_training_data_subset  
  IF Y numerical:  
    output <- average_Y_value_of_points_in_training_data_subset  
  return output  
END FUNCTION
```

Your answer is correct.

Correct

Marks for this submission: 4/4.

Question 7

Correct

Mark 4 out of 4

Lastly, fill in the comments that describe the steps for finally pruning the tree. The following function is the one that returns for a given subtree its pruned version (i.e., the original subtree, if no pruning should be applied, or else its pruned version).

```
FUNCTION prune_subtree(DecisionTree subtree, ValidationSet v):
  root <- root node of subtree

  # 1. 
  root.v_samples <- samples in v that lie in the subspace defined by this root node
  root.output <- determine_output(root.v_samples)
  mapping_without_splitting <- function g on root.v_samples: g(point) = root.output
  mapping_with_splitting <- function g' on root.v_samples: g(point) = leaf.output for leaf in subtree with point in leaf's subspace

  # 2. 
  performance_without_splitting <- performance indicator on root.v_samples of mapping_without_splitting
  performance_with_splitting <- performance indicator on root.v_samples of mapping_with_splitting

  # 3. 
  IF performance_with_splitting similar to performance_without_splitting:
    RETURN Leaf(root.set)
  ELSE:
    RETURN subtree
END FUNCTION
```

Your answer is correct.

Correct

Marks for this submission: 4/4.

Question 8

Correct

Mark 4 out of 4

Assume we are building a decision tree for binary classification into blue, pink. We have reached a decision tree node which only holds the sample set S of size $K = |S|$, and we would like to break it further down. As a candidate for splitting we consider the X feature, splitting at the split point x to get a left side (K_r samples in S , P_r of which are blue) and a right side (K_l samples in S , P_l of which are blue) subset.

Recap of the formulas

Let's calculate the decrease in impurity with respect to different purity measurements. For that, recall the formulas for this simple 2 input, 2 output class case:

- Information Gain [bits]:


$$I(X; Y) = H(Y) - H(Y | X) = -\left(\frac{P_r + P_l}{K} \log_2\left(\frac{P_r + P_l}{K}\right) + \left(1 - \frac{P_r + P_l}{K}\right) \log_2\left(1 - \frac{P_r + P_l}{K}\right)\right) + \left(\left(\frac{K_r}{K} \frac{P_r}{K_r} \log_2\left(\frac{P_r}{K_r}\right) \frac{K_r}{K} \left(1 - \frac{P_r}{K_r}\right) \log_2\left(1 - \frac{P_r}{K_r}\right)\right) + \left(\frac{K_l}{K} \frac{P_l}{K_l} \log_2\left(\frac{P_l}{K_l}\right) \frac{K_l}{K} \left(1 - \frac{P_l}{K_l}\right) \log_2\left(1 - \frac{P_l}{K_l}\right)\right)\right)$$

- Gini impurity decreases as

$$\text{Gini}(X = \text{any}) - (P(X \leq x) \text{Gini}(X \leq x) + P(X \geq x) \text{Gini}(X \geq x)) = \left(1 - \left(\left(\frac{P_r + P_l}{K}\right)^2 + \left(1 - \frac{P_r + P_l}{K}\right)^2\right)\right) - \left(\frac{K_r}{K} \left(1 - \left(\frac{P_r}{K_r}\right)^2 + \left(1 - \frac{P_r}{K_r}\right)^2\right) + \frac{K_l}{K} \left(1 - \left(\frac{P_l}{K_l}\right)^2 + \left(1 - \frac{P_l}{K_l}\right)^2\right)\right)$$

A concrete example

Recall the example from the lecture ($K_r = 10, P_r = 1, K_l = 8, P_l = 7$):



Let's manually go through the example using the formula for information gain in order to get a better gut feeling about their meaning.

- $H(Y) = -(P(\text{blue}) \log_2(P(\text{blue})) + P(\text{pink}) \log_2(P(\text{pink}))) = -(8/18) \log_2(8/18) - (10/18) \log_2(10/18) \approx 0.991$
- Split at $X_1 = x_1$ obtaining the new binary variable $X'_1 := (X_1 \geq x_1)$ (write $(X_1 \leq x_1) \Leftrightarrow$: left, $(X_1 > x_1) \Leftrightarrow$: right):
 - $H(Y | X'_1) = -P(\text{right}) (P(\text{blue} | \text{right}) \log_2(P(\text{blue} | \text{right})) + P(\text{pink} | \text{right}) \log_2(P(\text{pink} | \text{right}))) - P(\text{left}) (P(\text{blue} | \text{left}) \log_2(P(\text{blue} | \text{left})) + P(\text{pink} | \text{left}) \log_2(P(\text{pink} | \text{left})))$
 - $I(X'_1; Y) =$
- Split at $X_2 = x_2$ obtaining the new binary variable $X'_2 := (X_2 \geq x_2)$ (write $(X_2 \leq x_2) \Leftrightarrow$: bottom, $(X_2 > x_2) \Leftrightarrow$: top):
 - $H(Y | X'_2) = -P(\text{bottom}) (P(\text{blue} | \text{bottom}) \log_2(P(\text{blue} | \text{bottom})) + P(\text{pink} | \text{bottom}) \log_2(P(\text{pink} | \text{bottom}))) - P(\text{top}) (P(\text{blue} | \text{top}) \log_2(P(\text{blue} | \text{top})) + P(\text{pink} | \text{top}) \log_2(P(\text{pink} | \text{top})))$
 - $I(X'_2; Y) =$

This comparison clearly shows that splitting at feature X_1 is preferable, which is coherent with the visualization.

Correct

Marks for this submission: 4/4.

Question 9

Correct

Mark 4 out of 4

In the lecture we have visited bagging and boosting as means to build ensembles that reduce variance and bias compared to their single members. While the result of both is an ensemble, it is important to differentiate between them, because they follow quite different ideas and have different effects on the bias-variance tradeoff. Which of the following properties holds for which of the methods?

bagging	boosting	
<input checked="" type="radio"/>	<input type="radio"/>	Learning of ensemble members can be done in parallel.
<input checked="" type="radio"/>	<input type="radio"/>	Aims to decrease variance.
<input type="radio"/>	<input checked="" type="radio"/>	Aims to decrease bias.
<input type="radio"/>	<input checked="" type="radio"/>	Can be formulated to weight each training data point according to whether it is erroneously treated. Utilizes the complete data set (but weighted) for training each ensemble member.
<input type="radio"/>	<input checked="" type="radio"/>	Tries to circumvent instability of the learning method by active search for complementary solutions.
<input checked="" type="radio"/>	<input type="radio"/>	A special application to decision trees is random forests.
<input type="radio"/>	<input checked="" type="radio"/>	Learning of ensemble members is done sequentially.
<input checked="" type="radio"/>	<input type="radio"/>	Utilizes bootstrapping, i.e., resampling from a sample to obtain new sample sets. Every ensemble member is trained on a different sample set.
<input type="radio"/>	<input checked="" type="radio"/>	Builds a strong learner as an ensemble of weak learners.
<input checked="" type="radio"/>	<input type="radio"/>	Captures instability of the learning method by randomization.

Correct

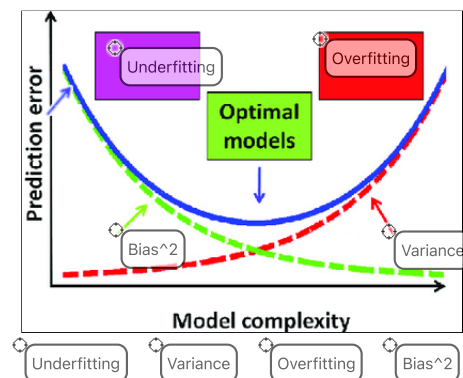
Marks for this submission: 4/4.

Question 10

Correct

Mark 4 out of 4

The bias, also the training error, and variance together make up the final testing error as depicted below. Correctly insert the right labels.



Your answer is correct.

Correct

Marks for this submission: 4/4.

Recap: Why is the bias-variance tradeoff really important?

Recall that one of the main learning goals in this semester is to be able to do informed model selection for your problem. For supervised models, this culminates in the bias-variance tradoff: The bias vs. variance tradeoff is one of the most important factors that should be taken into account for selecting the machine learning model to apply. Training a too complex/flexible model on training data for a simple model will almost for sure lead to overfitting, i.e., a high variance between solutions for different training sets; and on the other hand, selecting a too simple model for a complex problem cannot lead to a trustworthy solution.

Recap: How to influence bias and variance?

We have seen several ways how one can directly influence the **bias**:

- **By the choice of model:** fewer parameters respectively less *flexibility*, as we called it, means less bias. In particular, different models assume different shapes of problems that they work best for: Try to find the matching one.
- By adding further constraints to the model like regularization
- By accepting remaining model errors and simply fix it with error-fixing models, building an ensemble, which is boosting.

And for a large bias, the **variance** can only be reduced by

- decreasing the bias,
- random ensembling such as bagging, or
- by drastically increasing the amount of training samples.

Recap: Know your bias and variance!

To finalize the remark: Always make sure that you measure your bias and variance by applying testing on a separate test dataset to measure the generalization error, and cross-validation! Cross-validation will give you a glimpse on the variance of your model, while the generalization error for the average model gives you (by definition) an estimation for the bias.

Now let's finally come to the actual exercise: Try yourself on a couple of examples.

Rather high bias	Rather low bias		
<input checked="" type="radio"/>	<input type="radio"/>	Regression: Polynomial regression $f(x) = \theta_2 x^2 + \theta_1 x + \theta_0$	
<input checked="" type="radio"/>	<input type="radio"/>	Classification: k-NN with $k = 18$	
<input checked="" type="radio"/>	<input type="radio"/>	Regression: Linear regression model $f(x) = \theta^T x + b$	
<input type="radio"/>	<input checked="" type="radio"/>	Classification: k-NN with $\{(k=3)\}$	
<input type="radio"/>	<input checked="" type="radio"/>	Classification: Decision tree of depth 8	
<input type="radio"/>	<input checked="" type="radio"/>	Natural language processing: Llama 2 70B language model	This <i>definitely</i> has a minuscule bias, and variance is (again not measured and) only controlled by the immense amount of training data!
<input checked="" type="radio"/>	<input type="radio"/>	Classification: Quadratic discriminant analysis	
<input checked="" type="radio"/>	<input type="radio"/>	Regression/classification: Deep neural network with 1 hidden layer and 5 neurons.	
<input type="radio"/>	<input checked="" type="radio"/>	Regression: Polynomial regression $\{f(x) = \sum_{i=0}^7 \theta_i x^i\}$	
<input checked="" type="radio"/>	<input type="radio"/>	Image classification: AlexNet convolutional neural network for image classification on ImageNet classes	While AlexNet already features quite a couple of parameters, it still must be considered small and biased in the context of complex tasks like the classification of 1000 image classes.
<input checked="" type="radio"/>	<input type="radio"/>	Classification: Naive Bayes classifier	Note that this even has a higher bias than logistic regression, as it adds the conditional independence assumption on the inputs.
<input checked="" type="radio"/>	<input type="radio"/>	Classification: Logistic regression	
<input type="radio"/>	<input checked="" type="radio"/>	Object detection: YOLOv8 nano	Note that the creators of the YOLOv8 models do not provide variance measures; this is most probably due to the high training effort and the availability of sufficient data for proper cross-validation, not because the models have a low variance!

